

ShipShape Audit — Findings Summary

Phase 1 (Diagnosis) · branch shipshape/audit · 7/7 categories baselined · companion to AUDIT_REPORT.md · each table uses the PRD's prescribed metrics; each category includes **improvement options** (audit opportunities — no code changed)

Condition of record (all baselines): documents=577, issues=328, sprints=35, users=31 (pnpm db:seed + seed-augment.sql, deterministic). Node v20.20.2 · PostgreSQL 16 (Docker) · macOS aarch64. Every number is produced by a committed re-runnable script in scripts/audit/ → raw in docs/audit/raw/; Phase-2 re-runs the same scripts for before/after. **Data-safety:** pnpm --filter @ship/api test truncates ship_dev.

1 — Type Safety measured

TypeScript Compiler API AST walk (no linter in repo; strict ON). Scope A = non-test src. → cat1-type-safety.mjs

Total any types	94 (api 65 · web 29 · shared 0)
Total type assertions (as)	433 (api 135 · web 298 · shared 0)
Total non-null assertions (!)	325 (api 292 · web 33 · shared 0)
Total @ts-ignore / @ts-expect-error	0 (1 in tests only)
Strict mode enabled?	Yes (strict + noUncheckedIndexedAccess + noImplicitReturns + noFallthroughCasesInSwitch; web tsconfig omits last 3)
Strict mode error count (if disabled)	N/A (strict enabled)
Top 5 violation-dense files	weeks.ts (84) · projects.ts (51) · issues.ts (48) · UnifiedDocumentPage.tsx (37) · seed.ts (35)

Aggregate = 852 violations; PRD Improvement Target "eliminate 25%" ⇒ ≤ 639 (−213).

- High** — non-null ! (325, api 292) is the dominant risk class.
- High** — as (433, web 298): downstream loss of discriminated-union narrowing.
- Medium** — explicit any (94) in api routes; **no linter** guards new violations.

► **Improve** — PRD target: eliminate **25%** (852 → ≤639), real fixes only.

- A L-risk** remove non-null ! in api hot files (weeks 48 · issues 37 · seed 35 · team 28 = ~148) via real guards/early-returns.
- B M** replace web as casts (UnifiedDocumentPage/Editor/PropertiesPanel ~88) by narrowing on the document_type union.
- C L** add @typescript-eslint (no-non-null-assertion / no-explicit-any) as errors → durable guard.
- D M** type the y-protocols.d.ts shim + api-route any (~40) with real types.

Recommended: A + C — clears the count fast + stops regressions. Verify cat1-type-safety.mjs before/after.

2 — Bundle Size measured

Production vite build --sourcemap + sourcemap byte attribution. → cat2-bundle.mjs

Total production bundle size	2,275 KB raw / 695.5 KB gzip (JS 2,210 / CSS 65)
Largest chunk (name + size)	index-C2vAyoQ1.js — 2,025 KB raw / 575.7 KB gzip
Number of chunks	261 JS + 1 CSS — but 91.6% of all JS is in that one chunk
Top 3 largest dependencies	emoji-picker-react (~398 KB, 7.8%) · highlight.js (~376 KB, 7.4%) · react-router (~347 KB, 6.8%)
Unused dependencies identified	@tanstack/query-sync-storage-persister (candidate — verify)

Code splitting (PRD "How to Measure"): minimal — 3 React.lazy; tab chunks 1–16 KB; everything else monolithic. PRD Improvement Target = −15% total or −20% initial via splitting.

- High** — monolithic entry chunk (576 KB gzip, 91.6% of JS); no route/vendor split.
- High** — editor-only deps (highlight.js, emoji-picker) ship on first load; lazy-load = large low-risk win.

► **Improve** — PRD target: −15% total or −20% initial via splitting (no functionality removed).

- A L** lazy import() editor-only deps (emoji-picker ~398 KB + highlight.js ~376 KB) — load only when editor mounts (~25–35% initial).
- B L** highlight.js: register only used languages, not the full bundle (~300 KB raw).
- C M** route-level React.lazy+Suspense for top-level pages.
- D L** Vite manualChunks vendor split (react / tiptap+prosemirror / yjs) for caching.

Recommended: A + B — clears −20% initial at lowest risk. Verify cat2-bundle.mjs before/after.

3 — API Response Time measured

Dependency-free concurrent-load harness, API :3000 direct, warmup=10/budget=120/cell @ concurrency 10-25-50 (headline @25). → cat3-api.mjs

Endpoint (key flow) @conc 25	P50	P95	P99
GET /api/documents?document_type=wiki (main page, ~300 KB)	163ms	201ms	211ms
GET /api/issues (list issues, ~280 KB, 328 rows)	101ms	124ms	130ms
GET /api/weeks (sprint board)	20ms	25ms	27ms
GET /api/search/mentions?q=load (search)	16ms	20ms	21ms
GET /api/documents/:id (view document)	18ms	23ms	26ms

Concurrency tested 10/25/50, 0 errors all cells. P95 scales ~linearly on the slow two: main_page 100→201→**362ms**, list_issues 63→124→**222ms**.

1. **High** — two unbounded list endpoints dominate & degrade ~linearly; SQL <1.5 ms ⇒ cost = serializing 300/280 KB JSON on the shared event loop.
2. **Medium** — global rate limiter **100 req/min prod**; per-request session write on every endpoint.

► **Improve** — PRD target: **-20% P95** on ≥2 endpoints, identical conditions, root cause documented.

- **A M** drop heavy content/yjs_state blobs from the *list* responses of /api/documents & /api/issues (root cause = JSON serialization) → large P95 drop on both.
- **B M** add LIMIT + cursor pagination to the same two endpoints.
- **C L** enable compression middleware + cache headers on these GETs.
- **D H** stream/offload serialization off the shared event loop.

Recommended: A + C — A hits ≥2 endpoints at once; C is free margin. Verify cat3-api.mjs before/after.

4 — Database Query Efficiency measured

Postgres full statement logging; 5 marker-bracketed flows, API-pool PIDs only. → cat4-db.mjs

User Flow	Total Queries	Slowest (ms)	N+1?
Load main page /api/documents?type=wiki	4	1.419	No
View a document /api/documents/:id	4	0.214	No
List issues /api/issues	5	0.985	No
Load sprint board /api/weeks	5	0.227	No
Search content /api/search/mentions	5	0.433	No

EXPLAIN: main_page exec 0.144 ms / planning 0.561 ms; purpose-built partial index unused at this volume.

1. **High** — universal auth tax: every flow runs SELECT sessions + UPDATE sessions.last_activity — 2 of every 4-5 queries are auth overhead.
2. **Medium** — no pagination/LIMIT; grows linearly with workspace size. **Low** — well-indexed, no N+1.

► **Improve** — PRD target: **-20% query count** on ≥1 flow *or* -50% on slowest query (before/after EXPLAIN).

- **A L** throttle the per-request UPDATE sessions.last_activity (write only if >60 s stale) → view_document 4→3 = **-25%**, every flow loses 1 query.
- **B H** short-TTL session cache to also drop the per-request SELECT sessions (-50%; revocation-lag risk).
- **C M** collapse SELECT+UPDATE into one UPDATE ... RETURNING.
- **D M** pagination/LIMIT (shared w/ Cat-3 B) to bound rows scanned.

Recommended: A — lowest risk, exactly clears -20%, compounds Cat-3. Verify cat4-db.mjs before/after.

5 — Test Coverage & Quality

measured

Unit run 3× for flakiness; E2E static catalog; coverage attempted. → docs/audit/raw/cat5-before.txt

Total tests	451 unit (api) + 151 unit (web) + ~882 E2E = ~1,484 total (vs README "73+")
Pass / Fail / Flaky	api: 451 / 0 / 0 on fresh seed; 8 fail on the locked snapshot (test-isolation defect). web: 138 / 13 / — (TipTap/ProseMirror, hidden by default pnpm test). E2E: not obtainable.
Suite runtime	unit: ~15–16 s (16/15/16). E2E: not obtainable
Critical flows with zero coverage	document CRUD via HTTP, auth, real-time Yjs collaboration; 16 web unit files never run by pnpm test
Code coverage % (if measured)	api Lines 40.52% · Stmts 40.34% · Funcs 40.9% · Branches 33.44%. web Lines 28.53% · Stmts 27.63% · Funcs 25.6% · Branches 19.38%. Configured per PRD instruction (#3b: added @vitest/coverage-v8@4.0.17; absent as-shipped — finding kept).

1. **High** — mandated /e2e-test-runner skill **does not exist**; 882-test suite unrunnable by sanctioned method.
2. **High** — coverage unmeasurable as-shipped. **Positive** — unit layer stable & fast.

► **Improve** — PRD target: meaningful tests for 3 untested critical paths or fix 3 flaky w/ root cause (each test commented w/ risk it mitigates).

• **A M** add api integration tests for 3 untested paths: document CRUD over HTTP · auth/session lifecycle & timeout · Yjs collab persistence (locks in Cat-6 fix).

• **B L** isolated unit-test DB (separate DATABASE_URL) so tests stop truncating ship_dev — unblocks safe repeated runs.

• **C L** install @vitest/coverage-v8 + web coverage config → baseline measurable.

• **D H** implement the missing /e2e-test-runner skill → 882-test E2E runnable safely.

Recommended: B then A — B so tests don't destroy seed data, then A meets the bar. Verify pnpm --filter @ship/api test.

6 — Runtime Errors & Edge Cases

measured

Reproducible Playwright harness: console · malformed input · RT1 collab (validated control) · slow-3G · PG log. → cat6-runtime.mjs

Console errors during normal usage	0 across all 6 pages (0 console.error / pageerror / failed req / ≥500)
Unhandled promise rejections (server)	Not directly observable (no centralized API logging); proxy: 1 Postgres ERROR (the bad-UUID 500); 0 browser 5xx in normal nav
Network disconnect recovery	Pass — online control PERSISTED; transient disconnect (tab open) PERSISTED; Yjs IndexedDB offline store present
Missing error boundaries	None triggered in normal use (0 pageerrors); fault-injection not exercised this pass
Silent failures identified	text/plain body → 201 creates a default doc; + RT1/RT3 server-side swallowed persist

Malformed-input matrix: bad JSON → 400 HTML stack page · missing CSRF → 403 HTML stack · oversized title → 400 clean JSON · wrong content-type → 201 · bad UUID path → 500 + raw PG error

1. **High** — unvalidated path params → raw DB error as 500; bad-JSON/missing-CSRF leak HTML stack traces (inconsistent contract).
2. **Medium** — no Content-Type guard (silent doc creation). **RT1 honestly downgraded**: client offline path robust; residual = server-side swallowed persist.

► **Improve** — PRD target: fix 3 gaps; ≥1 must be real user-facing data-loss/confusion (repro + before/after + recording).

• **A M (data-loss)** surface the swallowed Yjs server-persist failure — stop the silent catch, signal client (toast+retry) & log; users currently believe failed saves succeeded.

• **B L** validate path params (UUID/zod) → **400/404** instead of raw PG error as 500.

• **C M** global JSON error envelope + prod handler → no HTML stack traces, consistent {error}.

• **D L** enforce Content-Type on writes → stop silent junk-doc creation (201).

Recommended: A + B + C — exactly 3; A is the mandatory data-loss fix. Verify cat6-runtime.mjs + repro.

7 — Accessibility measured

Lighthouse 13 a11y score/page — **3-run median**, JSON+HTML committed to `reports/a11y/before/` (per ShipShape LH guide) — + axe-core WCAG 2.1 A+AA + keyboard/contrast/landmark, validated by an unauthenticated login control. → `cat7-lighthouse.mjs` + `cat7-a11y.mjs`. Lighthouse/axe are automated (~30–40% of WCAG) and do NOT test keyboard traversal — high scores ≠ compliant.

Lighthouse accessibility score (per page)	median of 3 runs: login 98 · main_docs 91 · view_document 91 · issues 100 · my_week 96 · team_dir 100 (lowest = 91, <code>main_docs</code> & <code>view_document</code>). Full JSON+HTML in <code>reports/a11y/before/</code> . Plus axe violations/passes: login 0/23 · <code>main_docs</code> 2/22 · <code>view_document</code> 2/22 · issues 0/21 · <code>my_week</code> 1/20 · <code>team_dir</code> 0/20. <i>Automated scores miss keyboard — a 100 is not a victory.</i>
Total Critical/Serious violations	Critical = 2, Serious = 17 node instances (C+S = 19); 3 rules: <code>aria-required-children</code> , <code>listitem</code> , <code>color-contrast</code>
Keyboard navigation completeness	Broken (authenticated) — only 3 distinct focusable in 60 Tabs vs 4/4 login control, despite 369–1017 interactive els
Color contrast failures	15 nodes, all on <code>/my-week</code> (low-opacity muted text) — WCAG 1.4.3 AA
Missing ARIA labels or roles	<code>aria-required-children</code> (critical) + <code>listitem</code> (serious) on <code>main_docs</code> & <code>view_document</code> ; login lacks main/nav landmarks; <code>lang/title/single-h1</code> OK ✓

README "Section 508 / WCAG 2.1 AA compliant" claim → **CONTRADICTED** by the evidence above (automated scan = lower bound).

- High** — keyboard navigation broken on the authenticated app (WCAG 2.1.1/2.4.3, core 508).
- High** — README 508/WCAG-AA claim false as shipped. **Positive** — baseline hygiene solid; failures localized.

- **Improve** — PRD target: +10 Lighthouse on lowest page *or* fix **all Critical/Serious** on the 3 most important pages. (Lowest = `view_document` **91**; +10 = 101 > max 100, so the "fix all C/S" path is the realistic one.)
- A M** fix all Critical/Serious on `main_docs` · `view_document` · `my_week`: `aria-required-children`, `listitem` structure, 15 `color-contrast` nodes (proven by both `cat7-a11y` + `cat7-lighthouse`).
 - B H** fix authenticated-app keyboard nav (focus order / skip-link / tabbable interactive els) — the headline defect automated scores miss (3 → many reachable).
 - C L** contrast-only quick win: raise `.text-muted/*` token to clear 4.5:1 (removes 15 serious, +Lighthouse on `my_week`).
 - D M** add LHCI to CI so the Lighthouse metric stays measurable over time.

Recommended: A (matches the realistic PRD path, measurable via both Cat-7 harnesses) + **B** as the high-impact stretch.

Audit recommendation — top 3 highest-value opportunities (the per-category "► Improve" options above are the full set; these are the strongest, each maps to a PRD ≥20% target, is independently reproducible & low-blast-radius): **(1)** trim list-response payloads + paginate `/api/documents` & `/api/issues` → P95 drop on ≥2 endpoints (`cat3-api.mjs`). **(2)** throttle the per-request `last_activity` write → fewer queries on ≥1 flow (`cat4-db.mjs`). **(3)** lazy-load editor-only deps → smaller initial bundle (`cat2-bundle.mjs`).

Phase-1 gate: **7 / 7 categories baselined**; each table mirrors the PRD's prescribed metrics, and each category lists improvement **options** mapped to its PRD Improvement Target. These are audit *recommendations* (Phase-1 diagnostic output) — **no application code was changed**; only reproducible measurement instruments + deterministic test data. Full methodology, raw evidence, and the cross-category ranked findings live in `AUDIT_REPORT.md`. Each recommended fix would be executed and proven before/after with its named `scripts/audit/catN-*` harness under the fixed condition of record.